

RabbitMQ集群简介

- 第1章引言
 - 1.1 目的
 - 1.2 说明
 - 1.3 MQ
 - 1.4 概念
 - 1.5 MQ特点
 - 1.6 工作流程
 - 1.7 系统环境
- 第2章 RabbitMQ部署
 - 2.1 系统环境基本配置
 - 2.2 RabbitMQ配置
 - 2.3 RabbitMQ集群配置
- 第3章 RabbitMQ集群验证
 - 3.1 Nova配置MQ HA
 - 3.2 RabbitMQ HA验证
 - 3.3 RabbitMQ恢复
- 第4章 RabbitMQ知识普及
 - 4.1 RabbitMQ用户权限
 - 4.2 RabbitMQ集群
 - 4.3 RabbitMQ查询
 - 4.4 RabbitMQ其他
- 第5章 FAQ
 - 5.1 RabbitMQ集群-网络分区
- 第6章参考
 - 6.1 URL参考

第1章引言

1.1 目的

实现OpenStack RabbitMQ消息高可用性，实现MQ的负载均衡，缓解MQ压力，提高性能

1.2 说明

Rabbitmq的集群是依附于erlang的集群来工作的，所以必须先构建起erlang的集群景象。Erlang的集群中各节点是经由过程一个magic cookie来实现的，这个cookie存放在/var/lib/rabbitmq/.erlang.cookie中(本人采用的是yum安装)，文件是400的权限。所以必须包管各节点cookie对峙一致，不然节点之间就无法通信

1.3 MQ

MQ全称为Message Queue, 消息队列（MQ）是一种应用程序对应用程序的通信方法。应用程序通过读出入队列的消息（针对应用程序的数据）来通信，而无需专用连接来链接它们。消息传递指的是程序之间通过在消息中发送数据进行通信，而不是通过直接调用彼此来通信，直接调用通常是用于诸如远程过程调用的技术。排队指的是应用程序通过队列来通信。队列的使用除去了接收和发送应用程序同时执行的要求。其中较为成熟的MQ产品有IBM WEBSPPHERE MQ

1.4 概念

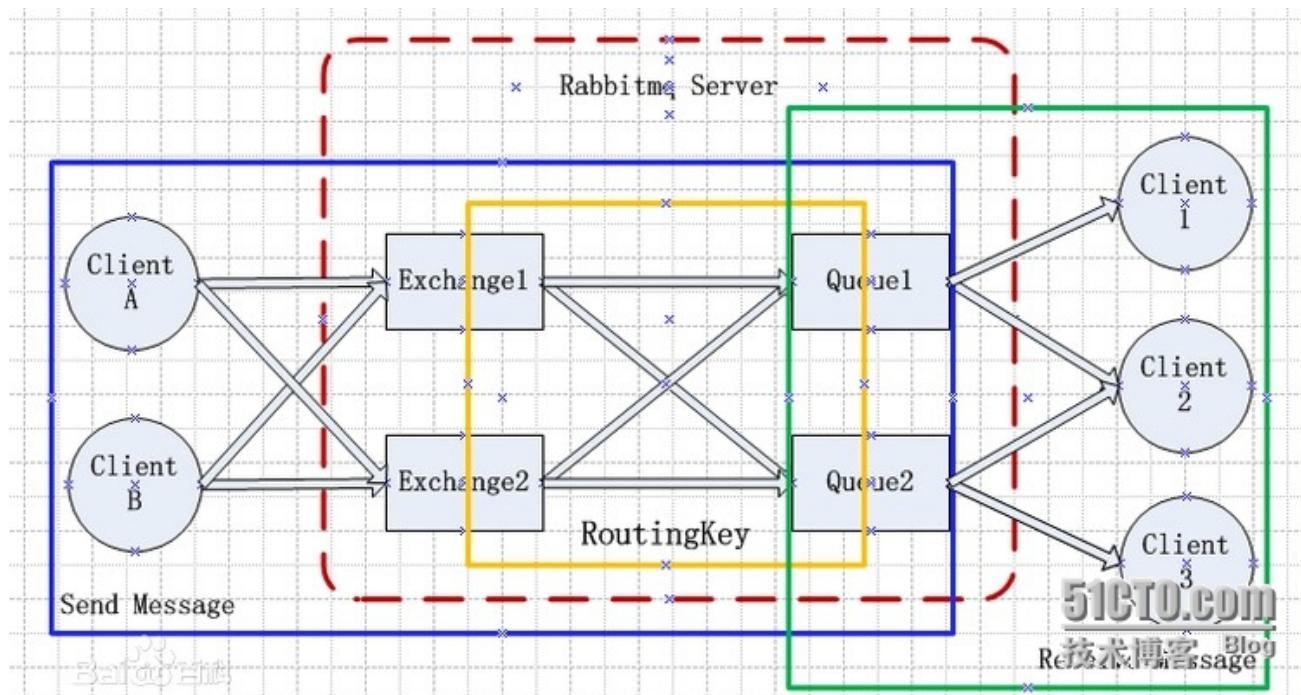
Broker：简单来说就是消息队列服务器实体。
Exchange：消息交换机，它指定消息按什么规则，路由到哪个队列。
Queue：消息队列载体，每个消息都会被投入到一个或多个队列。
Binding：绑定，它的作用就是把exchange和queue按照路由规则绑定起来。
Routing Key：路由关键字，exchange根据这个关键字进行消息投递。
vhost：虚拟主机，一个broker里可以开设多个vhost，用作不同用户的权限分离。
producer：消息生产者，就是投递消息的程序。
consumer：消息消费者，就是接受消息的程序。
channel：消息通道，在客户端的每个连接里，可建立多个channel，每个channel代表一个会话任务。

1.5 MQ特点

MQ是消费-生产者模型的一个典型的代表，一端往消息队列中不断写入消息，而另一端则可以读取或者订阅队列中的消息。MQ和JMS类似，但不同的是JMS是SUN JAVA消息中间件服务的一个标准和API定义，而MQ则是遵循了AMQP协议的具体实现和产品。

在项目中，将一些无需即时返回且耗时的操作提取出来，进行了异步处理，而这种异步处理的方式大大的节省了服务器的请求响应时间，从而提高了系统的吞吐量。

1.6 工作流程



消息队列的使用过程大概如下：

- (1) 客户端连接到消息队列服务器，打开一个channel。
- (2) 客户端声明一个exchange，并设置相关属性。
- (3) 客户端声明一个queue，并设置相关属性。
- (4) 客户端使用routing key，在exchange和queue之间建立好绑定关系。
- (5) 客户端投递消息到exchange。

Exchange接收到消息后，就根据消息的key和已经设置的binding，进行消息路由，将消息投递到一个或多个队列里。

Exchange也有几个类型，完全根据key进行投递的叫做Direct交换机，例如，绑定时设置了routing key为"abc"，那么客户端提交的消息，只有设置了key为"abc"的才会投递到队列。对key进行模式匹配后进行投递的叫做Topic交换机，符号"#"匹配一个或多个词，符号"*"匹配正好一个词。例如"abc.#"匹配"abc.def.ghi"，"abc.*"只匹配"abc.def"。还有一种不需要key的，叫做Fanout交换机，它采取广播模式，一个消息进来时，投递到与该交换机绑定的所有队列。

RabbitMQ支持消息的持久化，也就是数据写在磁盘上，为了数据安全考虑，我想大多数用户都会选择持久化。消息队列持久化包括3个部分：

- (1) exchange持久化，在声明时指定durable => 1
- (2) queue持久化，在声明时指定durable => 1
- (3) 消息持久化，在投递时指定delivery_mode=> 2 (1是非持久化)

如果exchange和queue都是持久化的，那么它们之间的binding也是持久化的。如果exchange和queue两者之间有一个持久化，一个非持久化，就不允许建立绑定。

1.7 系统环境

角色	主机名	IP地址	系统版本
MQ Node	rabbit@athController	192.168.8.180	CentOS6.5 + RabbitMQ3.1.5
MQ Node	rabbit@athBackup87	192.168.8.87	CentOS6.5 + RabbitMQ3.1.5

MQ Node	rabbit@athBackup53	192.168.8.53	CentOS6.5 + RabbitMQ3.1.5
---------	--------------------	--------------	---------------------------

第2章 RabbitMQ部署

本环境用于测试，采用三台节点作为MQ的集群，系统OS采用的是CentOS 6.5 x64, 该技术服务于OpenStack，卸载CentOS自带MQ产品QPID，采用的是RabbitMQ作为我们的首选高级消息队列服务

2.1 系统环境基本配置

- 1) install epel yum
 - a) `wget http://mirrors.yun-idc.com/epel/6/i386/epel-release-6-8.noarch.rpm`
 - b) `rpm -vi epel-release-6-8.noarch.rpm`
- 2) install rabbitMQ
 - a) `yum -y install rabbitmq-server`
- 3) Configure /etc/hosts （这步需要配置）
 - a) `192.168.8.180athController.8.180.abs.com.cnathController`
 - b) `192.168.8.53athBackup53.abs.com.cnathBackup53`
 - c) `192.168.8.87athBackup87.abs.com.cnathBackup87`
- 4) Configure /etc/sysconfig/network setting (Setting hostname, login other node Configure)
 - a) `sed -i 's/HOSTNAME=.*HOSTNAME=hostname/' /etc/sysconfig/network`
- 5) Disabled selinux Configure and iptables
 - a) `sed -i 's/SELINUX=.*SELINUX=disabled/' /etc/selinux/config`
 - b) `service iptables stop`
 - c) `chkconfig --level 123456 iptables off`
- 6) 以上操作，在三台节点同时执行操作

2.2 RabbitMQ配置

- 1) Configure /etc/rabbitmq/rabbitmq.config

```
[root@athController rabbitmq]# cat rabbitmq.config
% This file managed by Puppet

% Template Path: rabbitmq/templates/rabbitmq.config

[
  {rabbit, [
    {default_user, <<"guest">>},
    {default_pass, <<"guest">>}
  ]},
  {kernel, [
  ]}
].

% EOF

[root@athController rabbitmq]#
```

2) Configure/etc/rabbitmq/rabbitmq-env.config

```
[root@athController rabbitmq]# cat/opt/rabbitmq-env.conf
```

```
RABBITMQ_NODE_PORT=5672
```

```
[root@athController rabbitmq]#
```

3) Create .erlang.cookie, consistent

```
[root@athController rabbitmq]# vim /var/lib/rabbitmq/.erlang.cookie
```

```
YAGISQRAHKOF CZMWRFMT
```

```
[root@athController rabbitmq]# chmod 400 /var/lib/rabbitmq/.erlang.cookie
```

```
[root@athController rabbitmq]# chown rabbitmq:rabbitmq/var/lib/rabbitmq/.erlang.cookie
```

4) Restart rabbitmq-server

a) Service rabbitmq-server restart

5) Check RabbitMQ server and Add autostart

```
[root@athController nova]# netstat -tnpl | grep 5672
```

```
tcp0 0 0.0.0.0:15672 0.0.0.0:*LISTEN 26946/beam.smp
```

```
tcp0 0 0.0.0.0:5672 0.0.0.0:*LISTEN 26946/beam.smp
```

```
tcp0 0 0.0.0.0:5672 0.0.0.0:*LISTEN 26946/beam.smp
```

```
[root@athController nova]# chkconfig --level 123456 rabbitmq-server on
```

6) 三台节点同时配置并启动MQ服务

2.3 RabbitMQ集群配置

1) 主节点配置

```
[root@athController nova] rabbitmqctl stop_app
```

```
[root@athController nova] rabbitmqctl reset
```

```
[root@athController nova]
```

```
/usr/lib/rabbitmq/bin/rabbitmq-plugins enable rabbitmq_management
```

```
[root@athController nova]
```

```
/usr/lib/rabbitmq/bin/rabbitmq-plugins enable rabbitmq_management_agent
```

```
[root@athController nova] rabbitmqctl start_app
```

2) 节点配置

```
[root@athController rabbitmq] rabbitmqctl stop_app
```

```
[root@athController rabbitmq] rabbitmqctl reset
```

```
[root@athController rabbitmq] rabbitmqctl start_app
```

```
[root@athControllerrabbitmq]
/usr/lib/rabbitmq/bin/rabbitmq-plugins enable rabbitmq_management
[root@athControllerrabbitmq]
/usr/lib/rabbitmq/bin/rabbitmq-plugins enable rabbitmq_management_agent
[root@athController rabbitmq] rabbitmqctl join_cluster--ram athController
[root@athControllernova] rabbitmqctl start_app
```

3) 添加RabbitMQ用户(三台机器同时配置)

```
rabbitmqctl add_user username guest
```

```
rabbitmqctl change_password guest guest
```

4) RabbitMQ镜像设置

```
rabbitmqctl set_policy ha-all "^" '{"ha-mode":"all"}'
```

3) 浏览RabbitMQ WEB监控



[Overview](#)
[Connections](#)
[Channels](#)
[Exchanges](#)
[Queues](#)
[Admin](#)

Overview

► Totals

▼ Nodes

Name	File descriptors (?)	Socket descriptors (?)	Erlang processes	Memory	Disk space	Uptime	Type
rabbit@athBackup53	22 1024 available	1 829 available	414 1048576 available	34.0MB 1.5GB high watermark 653.7MB low watermark	46.2GB	1h 45m	RAM
rabbit@athBackup87	111 65536 available	84 58890 available	1261 1048576 available	77.3MB 25.2GB high watermark 653.7MB low watermark	45.2GB	2h 11m	Disc Stats
rabbit@athController	70 1024 available	48 829 available	909 1048576 available	42.0MB 50.4GB high watermark 653.7MB low watermark	44.6GB	1h 46m	Disc *

► Ports and contexts

▼ Import / export definitions

Export

Filename for download:

`rabbit_athController`

[Download broker definitions](#) (?)

Import

Definitions file:

[选择文件](#) [未选择文件](#)

[Upload broker definitions](#) (?)

51CTO.com
技术博客 Blog

Web contexts

Context	Node	Bound to	Port	SSL	Path
RabbitMQ Management	rabbit@athBackup53	0.0.0.0	15672	◇	/
RabbitMQ Management	rabbit@athBackup87	0.0.0.0	15672	◇	/
RabbitMQ Management	rabbit@athController	0.0.0.0	15672	◇	/
Redirect to port 15672	rabbit@athBackup53	0.0.0.0	55672	◇	/
Redirect to port 15672	rabbit@athBackup87	0.0.0.0	55672	◇	/
Redirect to port 15672	rabbit@athController	0.0.0.0	55672	◇	/

51CTO.com
技术博客 Blog

[Overview](#) [Connections](#) [Channels](#) [Exchanges](#) [Queues](#) [Admin](#)

Queues

All queues

Filter:

110 items (show at most 1000)

Overview						Messages			Message	
Name	Node	Exclusive	Parameters	Policy	Status	Ready	Unacked	Total	incoming	delivery
ceilometer.agent.central	rabbit@athBackup87+2			ha-all	Idle					
ceilometer.agent.central.athController.8.180.autohome.com.cn	rabbit@athBackup87+2			ha-all	Idle					
ceilometer.agent.central_fanout_e76416c4700b4a84a3a2dae1bc06df52	rabbit@athController+2		AD	ha-all	Idle					
ceilometer.agent.compute	rabbit@athBackup87+2			ha-all	Idle					
ceilometer.agent.compute.athCompute.8.179.autohome.com.cn	rabbit@athBackup87+2			ha-all	Idle					
ceilometer.agent.compute.athController.8.180.autohome.com.cn	rabbit@athBackup87+2			ha-all	Idle					
ceilometer.agent.compute_fanout_c3d921fe04464c258b813a17bee06515	rabbit@athController+2		AD	ha-all	Idle					
ceilometer.agent.compute_fanout_ddcaa028ec9d42e184b7850f2f8ac0b7	rabbit@athController+2		AD	ha-all	Idle					

51CTO.com
技术博客 Blog[Overview](#) [Connections](#) [Channels](#) [Exchanges](#) [Queues](#) [Admin](#)

Connections

Filter:

130 items (only showing first 100)

Network						Overview		
Name	Protocol	Client	Node	From client	To client	Timeout	Channels	User name
192.168.8.179:60272	AMQP 0-8		rabbit@athController	0B/s (2.7kB total)	0B/s (106.5kB total)		1	guest
192.168.8.179:60273	AMQP 0-8		rabbit@athController	0B/s (805B total)	0B/s (535B total)		1	guest
192.168.8.179:60274	AMQP 0-8		rabbit@athController	8B/s (74.4kB total)	205B/s (5.2MB total)		1	guest
192.168.8.179:60275	AMQP 0-8		rabbit@athController	0B/s (819B total)	0B/s (538B total)		1	guest
192.168.8.179:60276	AMQP 0-8		rabbit@athController	0B/s (1.1kB total)	0B/s (589B total)		1	guest
192.168.8.179:60277	AMQP 0-8		rabbit@athController	0B/s (309.7kB total)	0B/s (7.0kB total)		1	guest
192.168.8.179:60278	AMQP 0-8		rabbit@athController	0B/s (1.1MB total)	0B/s (20.7kB total)		1	guest
192.168.8.179:60279	AMQP 0-8		rabbit@athController	265B/s (2.2MB total)	5B/s (43.2kB total)		1	guest
192.168.8.179:60280	AMQP 0-8		rabbit@athController	0B/s (2.7MB total)	0B/s (40.8kB total)		1	guest

51CTO.com
技术博客 Blog

192.168.8.180:38216	AMQP 0-8		rabbit@athBackup87	0B/s (326.3kB total)	0B/s (6.6kB total)		1	guest
192.168.8.180:40755	AMQP 0-8		rabbit@athController	0B/s (495B total)	0B/s (392B total)		1	guest
192.168.8.180:40756	AMQP 0-8		rabbit@athController	0B/s (495B total)	0B/s (392B total)		1	guest
192.168.8.180:40757	AMQP 0-8		rabbit@athController	0B/s (789B total)	0B/s (532B total)		1	guest
192.168.8.180:40758	AMQP 0-8		rabbit@athController	0B/s (831B total)	0B/s (541B total)		1	guest
192.168.8.180:40759	AMQP 0-8		rabbit@athController	0B/s (817B total)	0B/s (538B total)		1	guest
192.168.8.180:40760	AMQP 0-8		rabbit@athController	0B/s (495B total)	0B/s (392B total)		1	guest
192.168.8.180:40761	AMQP 0-8		rabbit@athController	0B/s (2.8kB total)	0B/s (107.6kB total)		1	guest
192.168.8.180:40762	AMQP 0-8		rabbit@athController	0B/s (495B total)	0B/s (392B total)		1	guest

51CTO.com
技术博客 Blog

第3章 RabbitMQ集群验证

3.1 Nova配置MQ HA

```
1) NovaRabbitMQ配置

[root@athControllernova]# cat /etc/nova/nova.conf|grep rabbit

rabbit_host=192.168.8.180

rabbit_port=5672

rabbit_hosts=192.168.8.180:5672, 192.168.8.53:5672, 192.168.8.87:5672

rabbit_use_ssl=False

rabbit_userid=guest

rabbit_password=guest

rabbit_virtual_host=/

rabbit_retry_interval=1

rabbit_retry_backoff=2

rabbit_max_retries=0

rabbit_ha_queues=True

[root@athController nova]#
```

3.2 RabbitMQ HA验证

- 2) stop RabbitMQ集群
 - a) [root@athController nova]#/etc/init.d/rabbitmq-server stop
 - b) Stopping rabbitmq-server: rabbitmq-server.
 - c) [root@athController nova]#

3) RabbitMQ监控报警

Global counts (?)

Connections: 84 Channels: 84 Exchanges: 17 Queues: 74 Consumers: 124

▼ Nodes

Name	File descriptors (?)	Socket descriptors (?)	Erlang processes	Memory	Disk space	Uptime	Type
rabbit@athBackup53	27 1024 available	2 829 available	356 1048576 available	33.4MB 1.5GB high watermark	46.2GB 63.7MB low watermark	2h 15m	RAM *
rabbit@athBackup87	110 65536 available	84 58890 available	1188 1048576 available	73.5MB 25.2GB high watermark	45.2GB 63.7MB low watermark	2h 41m	Disc Stats
rabbit@athController	Node not running						Disc

4) Nova错误日志

```

2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/kombu/syn.py", line 14, in blocking
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common return _sync_current(fun, *args, **kwargs)
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/kombu/syn.py", line 40, in _blocking_
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common return spawn(fun, *args, **kwargs).wait()
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/eventlet/greenthread.py", line 168, in wait
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common return self._exit_event.wait()
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/eventlet/event.py", line 116, in wait
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common return hubs.get_hub().switch()
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/eventlet/hubs/hub.py", line 187, in switch
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common return self.greenlet.switch()
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/eventlet/greenthread.py", line 194, in main
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common result = function(*args, **kwargs)
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/amqp/lib/client_0_8/channel.py", line 839, in exchange_declare
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common self._send_method((40, 10), args)
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/amqp/lib/client_0_8/abstract_channel.py", line 70, in _send_method
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common method_sig, args, content)
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/amqp/lib/client_0_8/method_framing.py", line 233, in write_method
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common self._dest.write_frame(1, channel, payload)
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/amqp/lib/client_0_8/transport.py", line 125, in write_frame
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common frame_type, channel, size, payload, 0xce))
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/eventlet/greenio.py", line 309, in sendall
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common tail = self.send(data, flags)
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/eventlet/greenio.py", line 295, in send
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common total_sent += fd.send(data[total_sent:], flags)
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common error: [Errno 104] Connection reset by peer
2014-08-06 16:40:44.112 26605 TRACE nova.openstack.common.rpc.common
2014-08-06 16:40:44.114 26605 INFO nova.openstack.common.rpc.common [-] Reconnecting to AMQP server on 192.168.8.180:5672
2014-08-06 16:40:44.115 26605 ERROR nova.openstack.common.rpc.common [-] AMQP server on 192.168.8.180:5672 is unreachable: [Errno 111] ECONNREFUSED. Trying again in 1 seconds.
2014-08-06 16:40:44.115 26605 INFO nova.openstack.common.rpc.common [-] Reconnecting to AMQP server on 192.168.8.180:5672
2014-08-06 16:40:45.157 26605 INFO nova.openstack.common.rpc.common [-] Connected to AMQP server on 192.168.8.180:5672
2014-08-06 16:41:14.828 26605 AUDIT nova.compute.resource_tracker [-] Auditing locally available compute resources
2014-08-06 16:41:15.458 26605 AUDIT nova.compute.resource_tracker [-] Free ram (MB): 112117
2014-08-06 16:41:15.458 26605 AUDIT nova.compute.resource_tracker [-] Free disk (GB): 483
2014-08-06 16:41:15.458 26605 AUDIT nova.compute.resource_tracker [-] Free VCPUS: 24
2014-08-06 16:41:15.458 26605 INFO nova.compute.resource_tracker [-] Compute_service record updated for athController.8.180.autohome.com.cn:athController.8.180.autohome.com.cn

```



5) 192.168.8.53: 5672接管消息队列，已正常工作

```

2014-08-06 16:43:16.042 26605 AUDIT nova.compute.resource_tracker [-] Auditing locally available compute resources
2014-08-06 16:43:16.670 26605 AUDIT nova.compute.resource_tracker [-] Free ram (MB): 112117
2014-08-06 16:43:16.670 26605 AUDIT nova.compute.resource_tracker [-] Free disk (GB): 483
2014-08-06 16:43:16.670 26605 AUDIT nova.compute.resource_tracker [-] Free VCPUS: 24
2014-08-06 16:43:16.746 26605 INFO nova.compute.resource_tracker [-] Compute_service record updated for athController.8.180.autohome.com.cn:athController.8.180.autohome.com.cn
2014-08-06 16:45:17.295 26605 AUDIT nova.compute.resource_tracker [-] Auditing locally available compute resources
2014-08-06 16:45:17.928 26605 AUDIT nova.compute.resource_tracker [-] Free ram (MB): 112117
2014-08-06 16:45:17.928 26605 AUDIT nova.compute.resource_tracker [-] Free disk (GB): 483
2014-08-06 16:45:17.928 26605 AUDIT nova.compute.resource_tracker [-] Free VCPUS: 24
2014-08-06 16:45:18.000 26605 INFO nova.compute.resource_tracker [-] Compute_service record updated for athController.8.180.autohome.com.cn:athController.8.180.autohome.com.cn

```



6) 依次停止8.53:5672消息服务，8.87:5672正常接管

```

2014-08-06 16:47:34.405 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/amqp/lib/client_0_8/method_framing.py", line 233, in write_method
2014-08-06 16:47:34.405 26605 TRACE nova.openstack.common.rpc.common self._dest.write_frame(1, channel, payload)
2014-08-06 16:47:34.405 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/amqp/lib/client_0_8/transport.py", line 125, in write_frame
2014-08-06 16:47:34.405 26605 TRACE nova.openstack.common.rpc.common frame_type, channel, size, payload, 0xce))
2014-08-06 16:47:34.405 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/eventlet/greenio.py", line 309, in sendall
2014-08-06 16:47:34.405 26605 TRACE nova.openstack.common.rpc.common tail = self.send(data, flags)
2014-08-06 16:47:34.405 26605 TRACE nova.openstack.common.rpc.common File "/usr/lib/python2.6/site-packages/eventlet/greenio.py", line 295, in send
2014-08-06 16:47:34.405 26605 TRACE nova.openstack.common.rpc.common total_sent += fd.send(data[total_sent:], flags)
2014-08-06 16:47:34.405 26605 TRACE nova.openstack.common.rpc.common error: [Errno 104] Connection reset by peer
2014-08-06 16:47:34.405 26605 TRACE nova.openstack.common.rpc.common
2014-08-06 16:47:34.406 26605 INFO nova.openstack.common.rpc.common [-] Reconnecting to AMQP server on 192.168.8.180:5672
2014-08-06 16:47:34.407 26605 ERROR nova.openstack.common.rpc.common [-] AMQP server on 192.168.8.180:5672 is unreachable: [Errno 111] ECONNREFUSED. Trying again in 1 seconds.
2014-08-06 16:47:35.407 26605 INFO nova.openstack.common.rpc.common [-] Reconnecting to AMQP server on 192.168.8.180:5672
2014-08-06 16:47:35.408 26605 ERROR nova.openstack.common.rpc.common [-] AMQP server on 192.168.8.180:5672 is unreachable: [Errno 111] ECONNREFUSED. Trying again in 3 seconds.
2014-08-06 16:47:38.410 26605 INFO nova.openstack.common.rpc.common [-] Reconnecting to AMQP server on 192.168.8.87:5672
2014-08-06 16:47:38.454 26605 INFO nova.openstack.common.rpc.common [-] Connected to AMQP server on 192.168.8.87:5672
2014-08-06 16:48:34.616 26605 AUDIT nova.compute.resource_tracker [-] Auditing locally available compute resources
2014-08-06 16:48:35.234 26605 AUDIT nova.compute.resource_tracker [-] Free ram (MB): 112117
2014-08-06 16:48:35.234 26605 AUDIT nova.compute.resource_tracker [-] Free disk (GB): 483
2014-08-06 16:48:35.235 26605 AUDIT nova.compute.resource_tracker [-] Free VCPUS: 24
2014-08-06 16:48:35.304 26605 INFO nova.compute.resource_tracker [-] Compute_service record updated for athController.8.180.autohome.com.cn:athController.8.180.autohome.com.cn

```



3.3 RabbitMQ恢复

1) 主节点恢复启动MQ

```
[root@athControllernova]# /etc/init.d/rabbitmq-server start
```

```
Starting rabbitmq-server: SUCCESS
```

```
rabbitmq-server.
```

```
[root@athController nova]#
```

2) RabbitMQ监控显示

Global counts (?)

Connections: 135

Channels: 135

Exchanges: 49

Queues: 110

Consumers: 248

▼ Nodes

Name	File descriptors (?)	Socket descriptors (?)	Erlang processes	Memory	Disk space	Uptime	Type
rabbit@athBackup53	22 1024 available	1 829 available	414 1048576 available	33.2MB 1.5GB high watermark	46.2GB 953.7MB low watermark	1m 42s	RAM
rabbit@athBackup87	112 65536 available	85 58890 available	1270 1048576 available	78.0MB 25.2GB high watermark	45.2GB 953.7MB low watermark	2h 59m	Disc
rabbit@athController	74 1024 available	52 829 available	945 1048576 available	39.8MB 50.4GB high watermark	44.6GB 953.7MB low watermark	1m 27s	

51CTO.com
技术博客 Blog



Overview Connections Channels Exchanges Queues Admin

Node rabbit@athBackup87

▼ Overview

File descriptors (?)	112 65536 available	Memory	78.3MB 25.2GB high watermark	Uptime	3h 1m
Socket descriptors (?)	85 58890 available	Disk space	45.2GB 953.7MB low watermark	Type	Disc
Erlang processes	1270 1048576 available				

▼ Memory details



Last updated: 2014-08-06 17:02:25.
Total memory used at last update: 78.5MB (?)

Update

51CTO.com
技术博客 Blog

第4章 RabbitMQ知识普及

4.1 RabbitMQ用户权限

1) 创建用户，配置密码

a) add_user <username><password>

2) 删除用户

a) delete_user <username>

3) 改变用户密码

a) change_password<username> <newpassword>

4) 清除用户密码

a) clear_password <username>

5) 设置用户权限

a) set_user_tags <username><tag>

b) rabbitmqctl set_user_tags rootadministrator

6) 查询用户权限

a) list_users

4.2 RabbitMQ集群

1) 添加RabbitMQ节点加入集群

a) join_cluster<clusternode> [--ram]

2) 查看集群状态

a) cluster_status

3) 修改集群node存储方式

a) change_cluster_node_type disc |ram

b) rabbitmqctl change_cluster_node_type ram

4) 跟新集群节点

a) update_cluster_nodes clusternode

5) RabbitMQ同步消息队列

a) sync_queue queue

b) rabbitmqctl sync_queue compute

4.3 RabbitMQ查询

1) 查询消息队列信息

a) list_queues [-p<vhostpath>] [<queueinfoitem> ...]

2) 查询消息交换机信息

a) list_exchanges [-p<vhostpath>] [<exchangeinfoitem> ...]

3) 查询exchanges队列直接bind信息

a) list_bindings [-p<vhostpath>] [<bindinginfoitem> ...]

4) 查询消息连接信息，如果是集群，将显示node bind信息

a) list_connections [<connectioninfoitem> ...]

4.4 RabbitMQ其他

1) RabbitMQ进程停止

a) stop [<pid_file>]

b) service rabbitmq-server

2) RabbitMQ应用停止与启动

a) stop_app, start_app

3) RabbitMQ重置/强制重置节点到原始状态

a) reset, force_reset

4) 查看RabbitMQ基础信息

5) Rabbitmqctl status

第5章 FAQ

5.1 RabbitMQ集群-网络分区

1) RabbitMQ集群网络分区问题

Overview

Network partition detected

Mnesia reports that this RabbitMQ cluster has experienced a network partition. This is a dangerous situation. RabbitMQ clusters should not be installed on networks which can experience partitions.

The nature of the partition is as follows:

Node	Was partitioned from
rabbit@athController	rabbit@athBackup53 rabbit@athBackup87

While running in this partitioned state, changes (such as queue or exchange declaration and binding) which take place in one partition will not be visible to other partition(s). Other behaviour is not guaranteed.

51CTO.com
技术博客 Blog

RabbitMQ clusters do not tolerate network partitions well. If you are thinking of clustering across a WAN, don't. You should use federation or the shovel instead.

RabbitMQ 一共有三种处理方式: ignore, autoheal, pause_minority. 默认的处理方式是 ignore, 即什么也不做

1) ignore 默认处理方式

```
[root@athController rabbitmq]# cat rabbitmq.config

% This file managed by Puppet

% Template Path: rabbitmq/templates/rabbitmq.config

[
  {rabbit, [
    {default_user, <<"guest">>},
    {default_pass, <<"guest">>},
    {tcp_listeners, [5672]},
    {cluster_partition_handling, ignore}
  ]},
  {kernel, [
  ]}
].

% EOF

[root@athController rabbitmq]#
```

2) autoheal 的处理方式: 简单来讲就是当网络分区恢复后, rabbitmq 各分区彼此进行协商, 分区中客户端连接数最多的为胜者, 其余的全部会进行重启, 这样也就恢复到同步的状态了

3) pause_minority 的处理方式: rabbitmq 节点感知集群中其他节点 down 掉时, 会判断自己在集群中处于多数派还是少数派, 也就是判断与自己形成集群的节点个数在整个集群中的比例是否超过一半。如果是多数派, 则正常工作, 如果是少数派, 则会停止 rabbit 应用并不断检测直到自己成为多数派的一员后再次启动 rabbit 应用。注意: 这种处理方式集群通常由奇数个节点组成。在 CAP 中, 优先保证了 CP

RabbitMQ 集群的网络分区容错性并不是非常高, 在网络经常发生分区时会有些问题, 最明显的就是脑裂问题, 官方给出的处理方式

Clustering and Network Partitions

RabbitMQ clusters do not tolerate network partitions well. If you are thinking of clustering across a WAN, don't. You should use [federation](#) or the [shovel](#) instead.

However, sometimes accidents happen. This page documents how to detect network partitions, some of the bad effects that may happen during partitions, and how to recover from them.

RabbitMQ stores information about queues, exchanges, bindings etc in Erlang's distributed database, Mnesia. Many of the details of what happens around network partitions are related to Mnesia's behaviour.

第6章参考

6.1 URL参考

<http://www.rabbitmq.com/ha.html>

<http://www.rabbitmq.com/shovel.html>

<http://www.rabbitmq.com/partitions.html>

<http://www.rabbitmq.com/partitions.html>

<http://www.rabbitmq.com/clustering.html>

<http://www.bbtang.info/linux/fuwu/610.html>

<http://www.rabbitmq.com/federation.html>

<http://baike.baidu.com/view/4095865.htm?fr=aladdin>

<http://www.kankanews.com/10kengine/archives/71918.shtml>

<http://lists.rabbitmq.com/pipermail/rabbitmq-discuss/2014-January/033412.html>

<http://rabbitmq.1065348.n5.nabble.com/Node-statistics-not-available-v3-2-0-td32937.html>

这是一篇网上的文章，为了重复造轮子，这里直接粘过来了，我们之后，可以根据自已的需求，在这个基础上进行修改。

连接地址：<http://swq499809608.blog.51cto.com/797714/1540205>